

## Lesson 13



# NoSQL Databases: Cassandra



# Database Problems

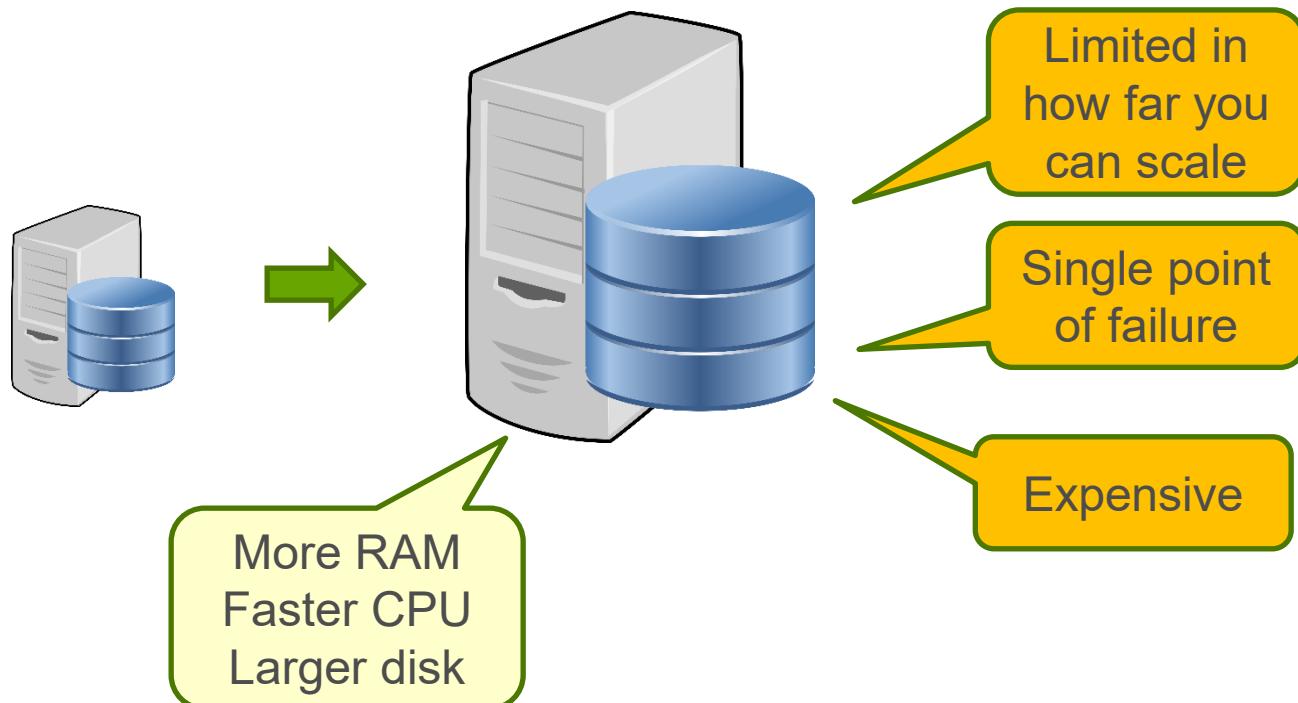
- Too much data
- Too much load on DB
  - DB becomes the bottleneck



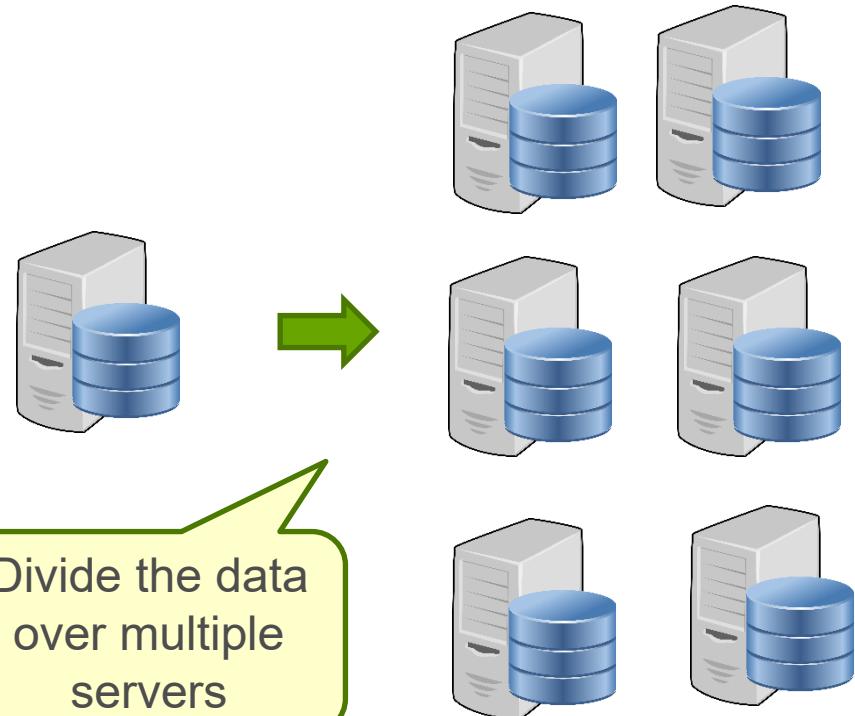


# Database Scaling

- Vertical scaling



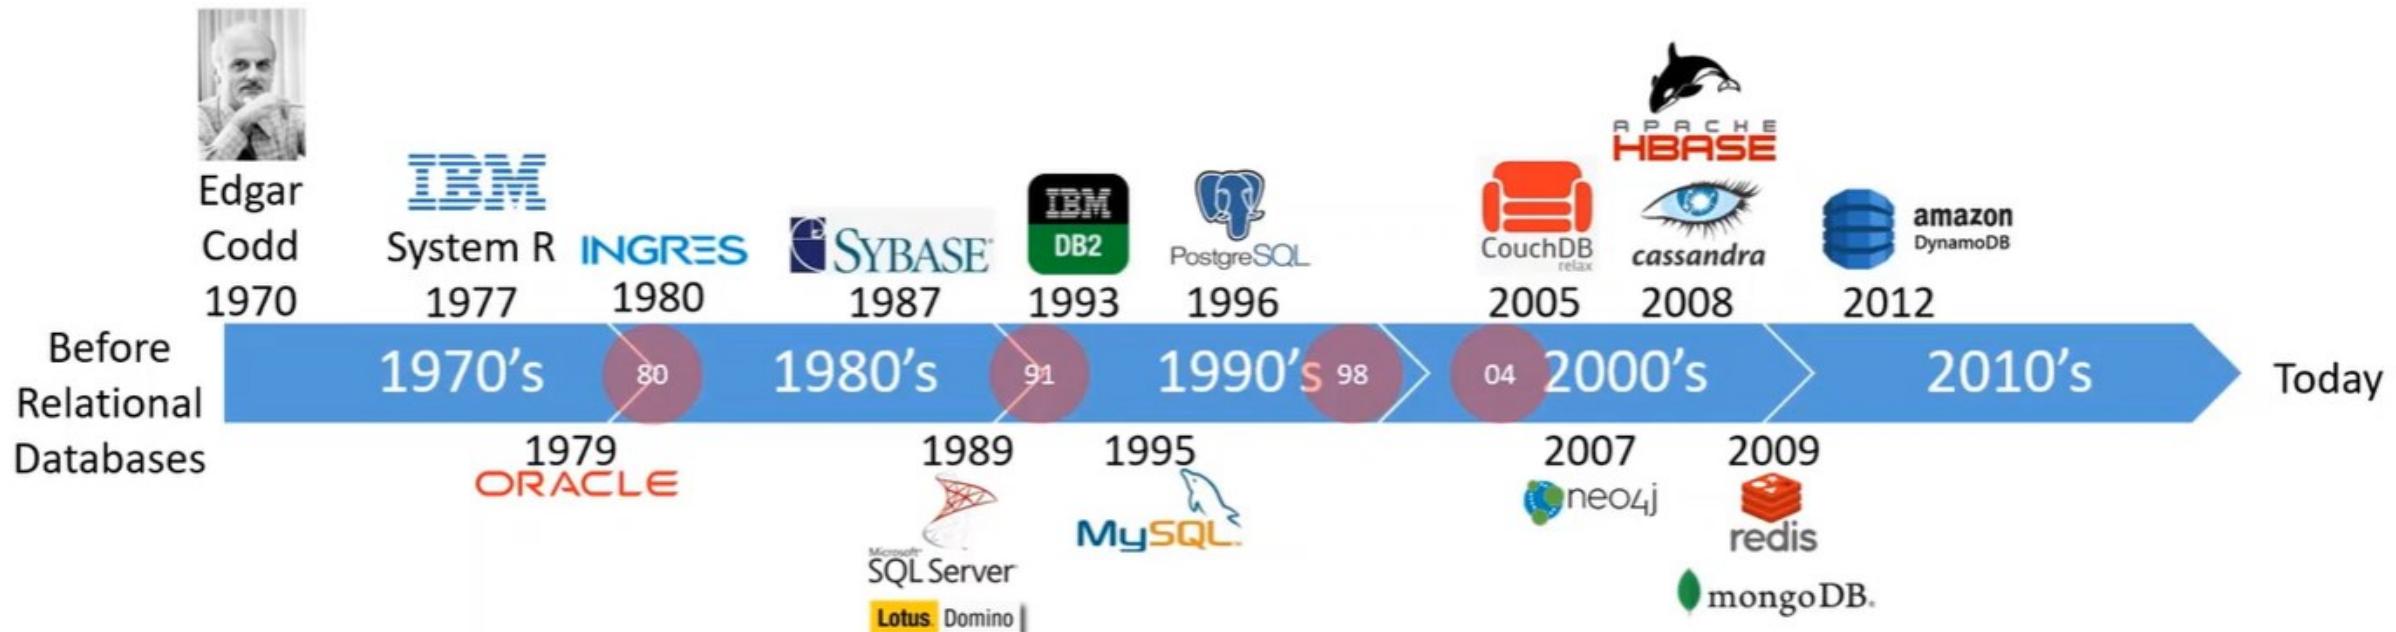
- Horizontal scaling





# Today's Requirements on Databases

- Big data (large datasets)
- Varied data formats
- Unstructured/ semi structured data





# Sharding in Distributed DBs

Original Table

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN
3	SELDÄ	BAĞCAN	PURPLE
4	JIM	PEPPER	AUBERGINE



HP1

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN

HP2

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
3	SELDÄ	BAĞCAN	PURPLE
4	JIM	PEPPER	AUBERGINE



# Key based Sharding

Shard  
Key

COLUMN 1	COLUMN 2	COLUMN 3
A		
B		
C		
D		

Hash values are created from the Shard Key

HASH  
FUNCTION



COLUMN 1	HASH VALUES
A	1
B	2
C	1
D	2

Shard 1

COLUMN 1	COLUMN 2	COLUMN 3
A		
C		

Shard 2

COLUMN 1	COLUMN 2	COLUMN 3
B		
D		



# Range-based Sharding

PRODUCT	PRICE
WIDGET	\$118
GIZMO	\$88
TRINKET	\$37
THINGAMAJIG	\$18
DOODAD	\$60
TCHOTCHKE	\$999



(\$0-\$49.99)

PRODUCT	PRICE
TRINKET	\$37
THINGAMAJIG	\$18



(\$50-\$99.99)

PRODUCT	PRICE
GIZMO	\$88
DOODAD	\$60



(\$100+)

PRODUCT	PRICE
WIDGET	\$118
TCHOTCHKE	\$999



# Problems with Relational Databases

- Scaling writes is very difficult and limited
  - Vertical scaling is limited and is expensive
  - Horizontal scaling is limited and is complex
    - Queries work only within shards
    - Strict consistency and partition tolerance leads to availability problems

*A relational database is hard to scale*



# Problems with Relational Databases contd..

- The schema in a database is fixed
- Schema evolution
  - Adding attributes to an object => have to add columns to table
  - You need to do a migration project
  - Application downtime ...

*A relational database is hard to change*



# Problems with Relational Databases contd..

- Relational schema doesn't easily handle unstructured and semi-structured data
  - Emails, Tweets, Pictures, Audio, Movies, Text

Unstructured data				Semi-structured data				Structured data					
ID	Name	Age	Degree	ID	Name	Age	Degree	ID	Name	Age	Degree		
1	John	18	B.Sc.	<University>	<Student ID="1">	<Name>John</Name>	<Age>18</Age>	<Degree>B.Sc.</Degree>	</Student>	2	David	31	Ph.D.
2	David	31	Ph.D.	<Student ID="2">	<Name>David</Name>	<Age>31</Age>	<Degree>Ph.D. </Degree>	</Student>	3	Robert	51	Ph.D.	
3	Robert	51	Ph.D.	....	</University>				4	Rick	26	M.Sc.	
4	Rick	26	M.Sc.						5	Michael	19	B.Sc.	
5	Michael	19	B.Sc.										

*A relational database does not handle unstructured and semi structured data very well*



# NoSQL Databases

- NoSQL databases, also known as "Not Only SQL" databases, have emerged in response to the limitations of traditional relational databases in handling modern, large-scale, and diverse types of data.
- The need for NoSQL databases arises from various challenges that traditional SQL databases face in certain scenarios.



# Advantages of NoSQL (Not Only SQL) Databases

- Capable to handle Big Data
- Suitable for Unstructured and Semi-Structured Data
- Horizontal Scalability
- Provide Flexible Schemas
- Optimized read and write Performance
- Provide Real-time Data processing
- High Availability and Fault Tolerance
- Cost-Effectiveness



# Choice between NoSQL and RDBMS

- It's important to note that while NoSQL databases offer many advantages, they are not one-size-fits-all solutions.
- The choice between NoSQL and traditional relational databases should be made based on the specific needs and characteristics of the application.
- RDBMS is a good choice for applications that require strong data consistency, integrity, and transactional support, and where data is well-structured and follows a defined schema.



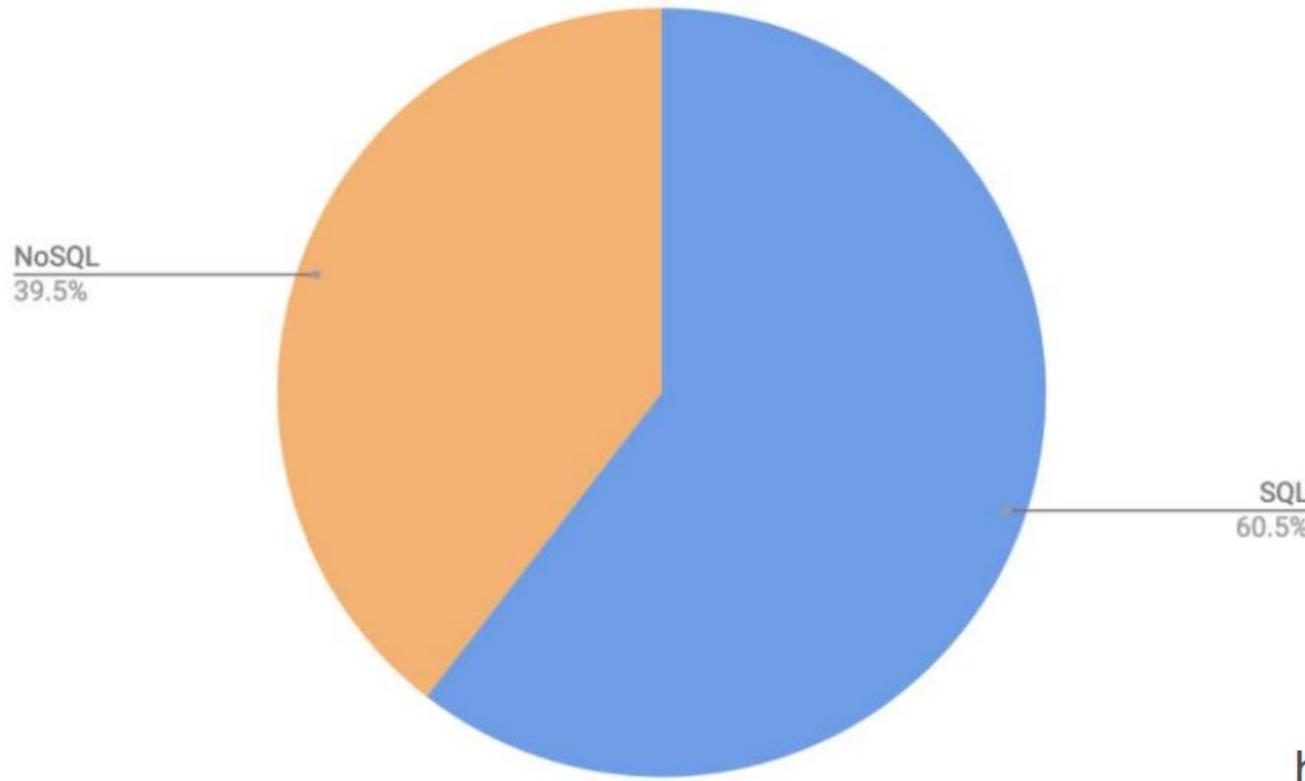
## Choice between NoSQL and RDBMS contd..

- NoSQL is a better choice for applications that need to handle large volumes of unstructured or semi-structured data, require extreme scalability, and can tolerate some level of eventual consistency.
- Many modern applications use a combination of both relational and NoSQL databases to meet their diverse data management requirements.



# Relational vs. No-SQL usage

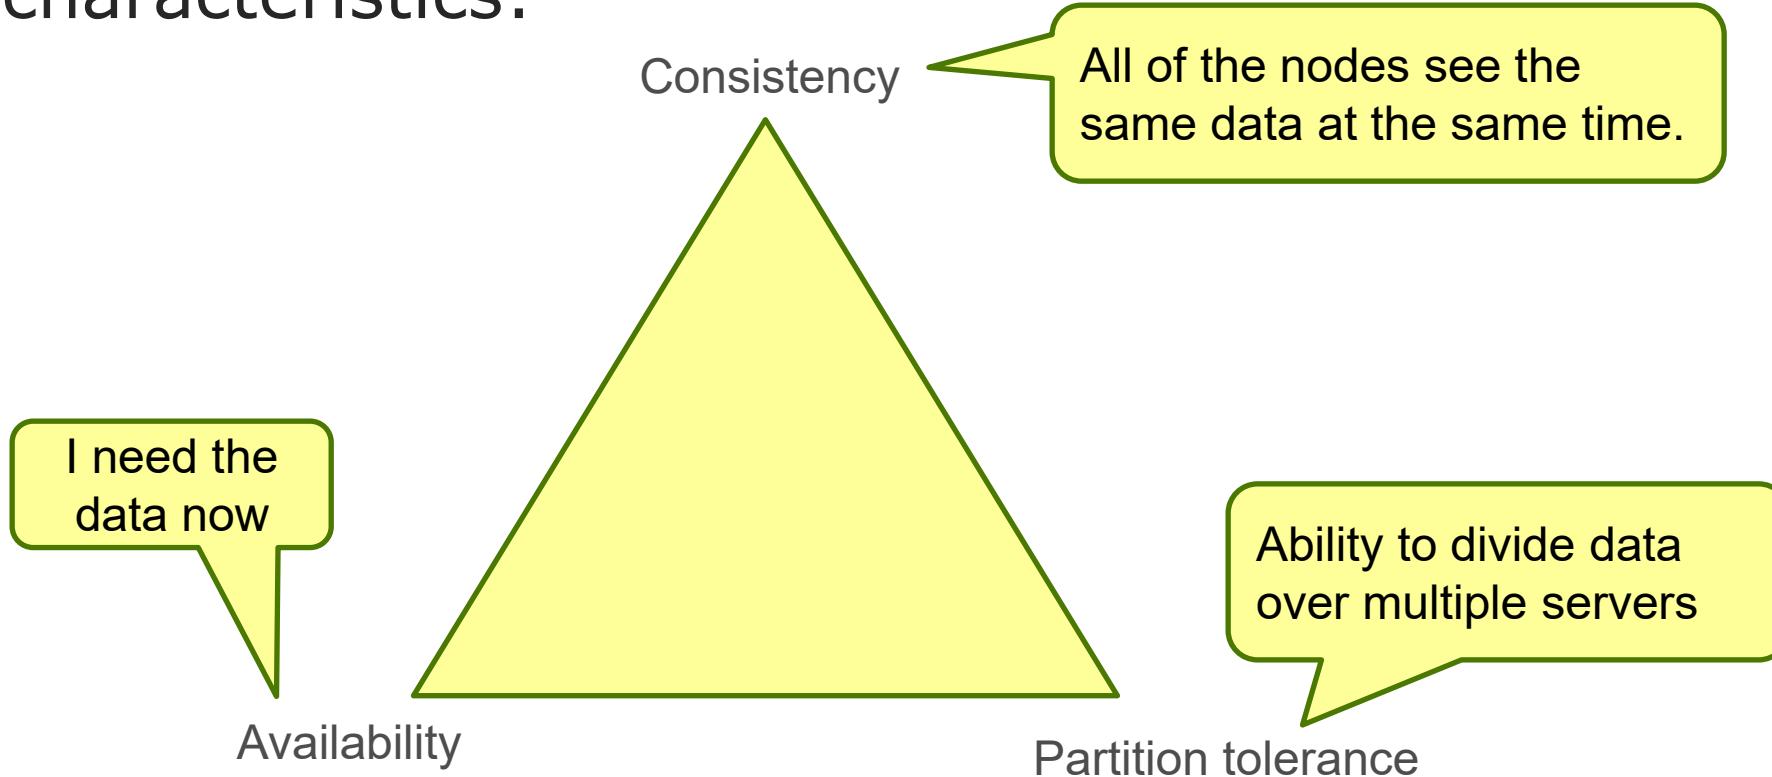
SQL Database Use: 60.48%  
NoSQL Database Use: 39.52%





# Brewer's CAP Theorem

- A distributed system can support only two of the following three characteristics:





# Consistency

- Strict consistency
  - The data that I read is always correct
  - You never loose data
- Eventual consistency
  - The data might not be correct
    - But will eventually become correct
- NoSQL databases prioritize flexibility and scalability over strict consistency.
- They are designed to handle diverse data types and large volumes.



# ACID - BASE

## ● Relational

### ● ACID

- **Atomicity:** All parts of transaction complete or none complete
- **Consistency:** Only valid data written to database
- **Isolation:** Parallel transactions do not impact each other's execution
- **Durability:** Once transaction committed, it remains

## ● NoSQL

### ● BASE

- Basically available
- Soft state
  - No schema
- Eventual consistency



# Types of NoSQL Databases

## Key Value



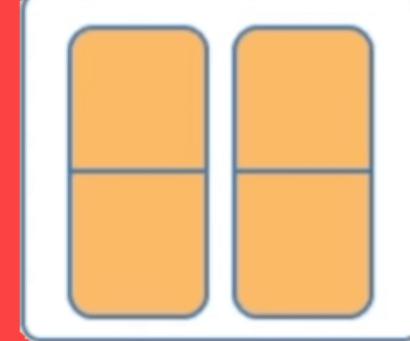
Example:  
Riak, Tokyo Cabinet, Redis  
server, Memcached,  
Scalaris

## Document-Based



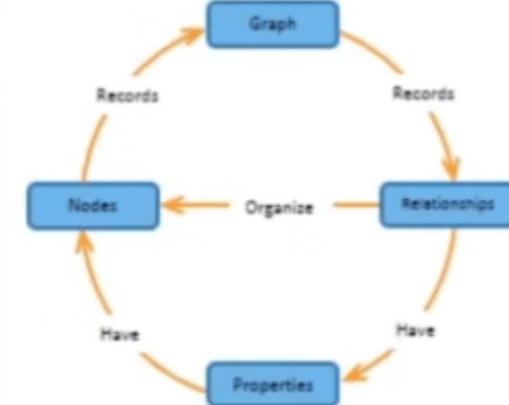
Example:  
MongoDB, CouchDB,  
OrientDB, RavenDB

## Column-Based



Example:  
BigTable, Cassandra,  
Hbase,  
Hypertable

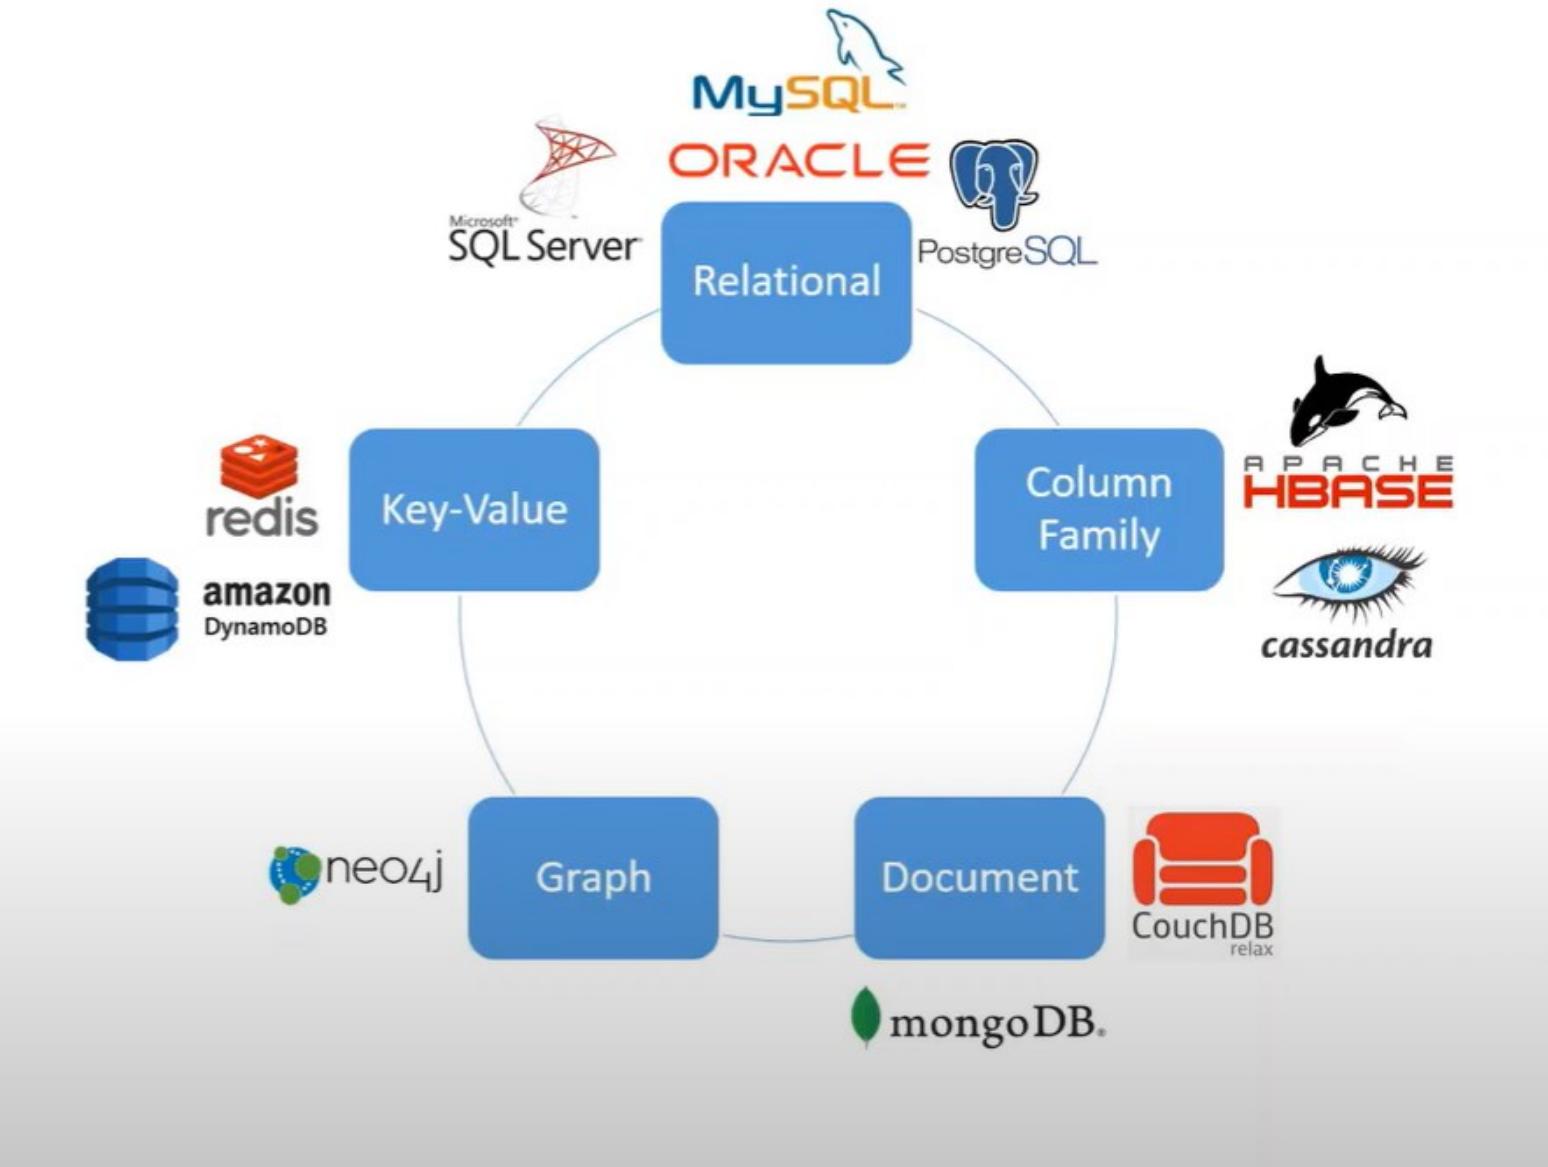
## Graph-Based



Example:  
Neo4J, InfoGrid, Infinite  
Graph, Flock DB



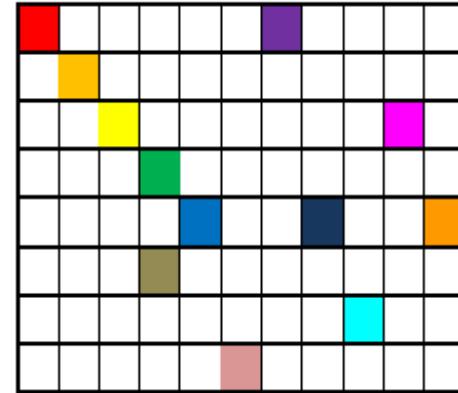
# Five Different Types of Databases





# Column family

- Place data in a certain column
- Ideal for high-variability data sets
- Column families allow to query all columns that have a specific property or properties
- Allow new columns to be inserted without doing an "alter table"

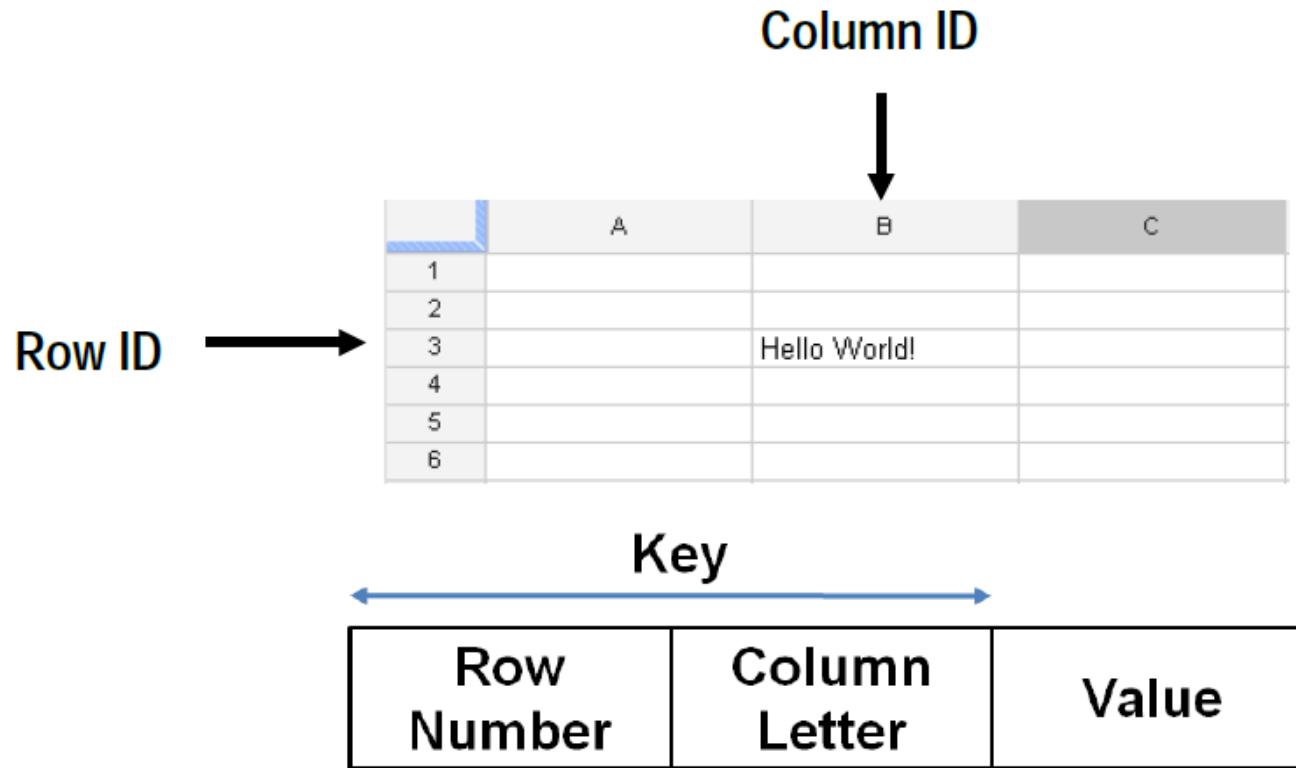


- Players
  - Cassandra
  - HBase
  - Google BigTable



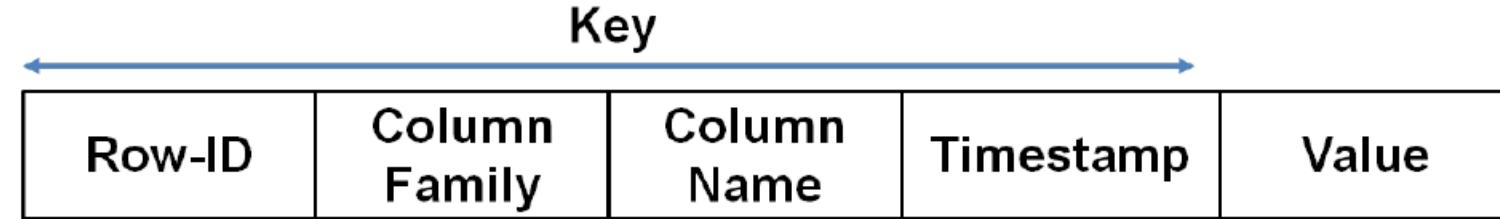
# Excel

- The key is the row number and the column letter





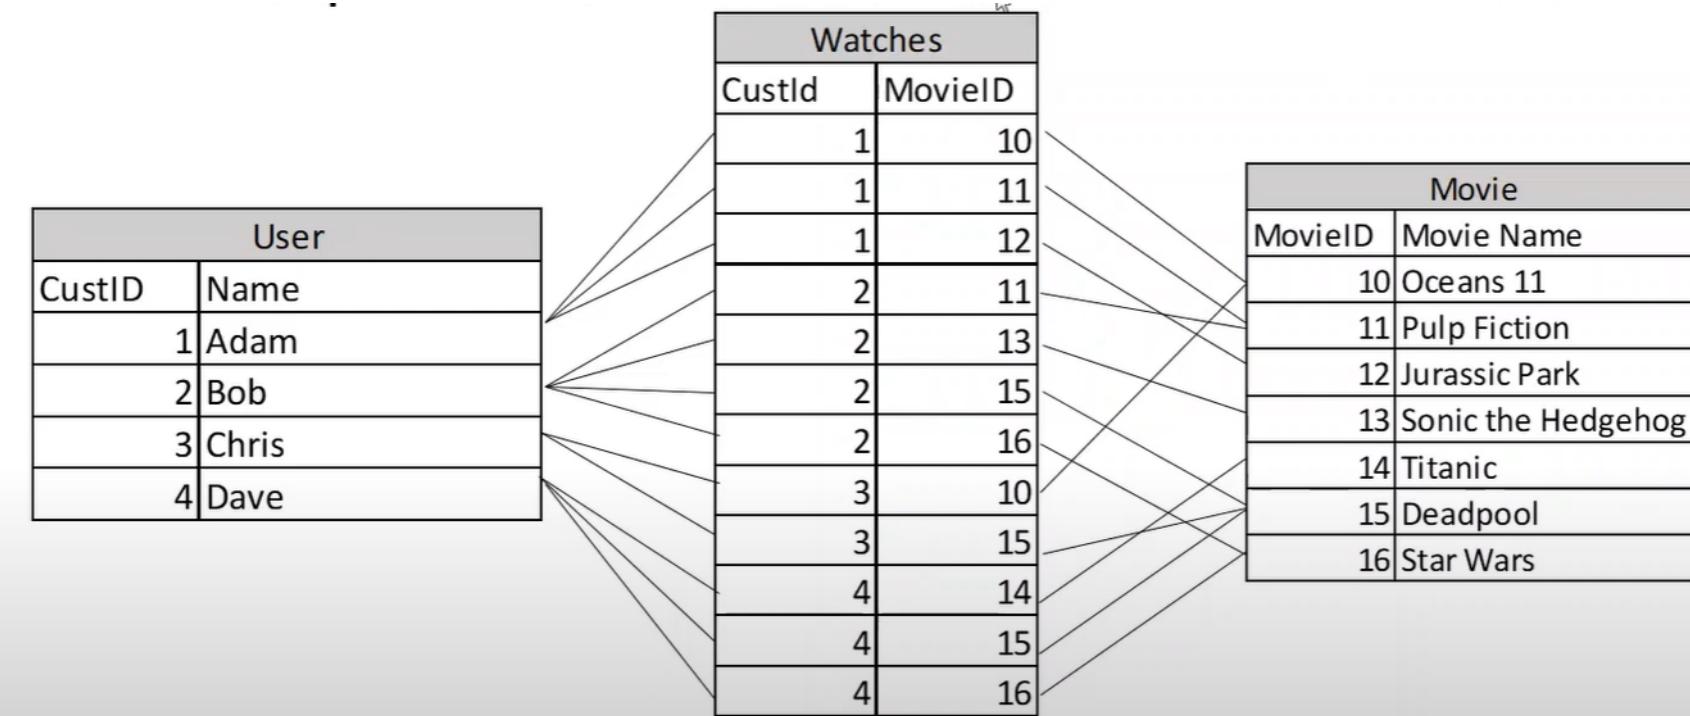
# Column family



- Keys can have many attributes
  - Column family (group columns in column families)
    - Group column families in super-families
  - Timestamp
- Queries can be done on rows, column families and column names, timestamp



# Relational



What movies has Adam watched?  
Who watched Deadpool?



Need to join 3 tables (slow)  
Need to join 3 tables (slow)



# Column family

User	Movies			
	Oceans 11	Pulp Fiction	Jurassic Park	
Adam	TRUE	TRUE	TRUE	
	Pulp Fiction	Sonic	Deadpool	Star Wars
Bob	TRUE	TRUE	TRUE	TRUE
	Oceans 11	Deadpool		
Chris	TRUE	TRUE		
	Titanic	Deadpool	Star Wars	
Dave	TRUE	TRUE	TRUE	

What movies has Adam watched?



Very easy and fast

Who watched Deadpool?



Difficult



# Column Family NoSQL Database: Cassandra

- Apache Cassandra was initially developed at Facebook and it came from Amazon's highly available Dynamo DB and Google's BigTable data model.
- It is a Column family NoSQL distributed database designed for scalability and high availability for Big Data.
- Unlimited elastically scalable
  - Installation is on multiple different nodes
- Always available (no downtime)
- Uses partitioning





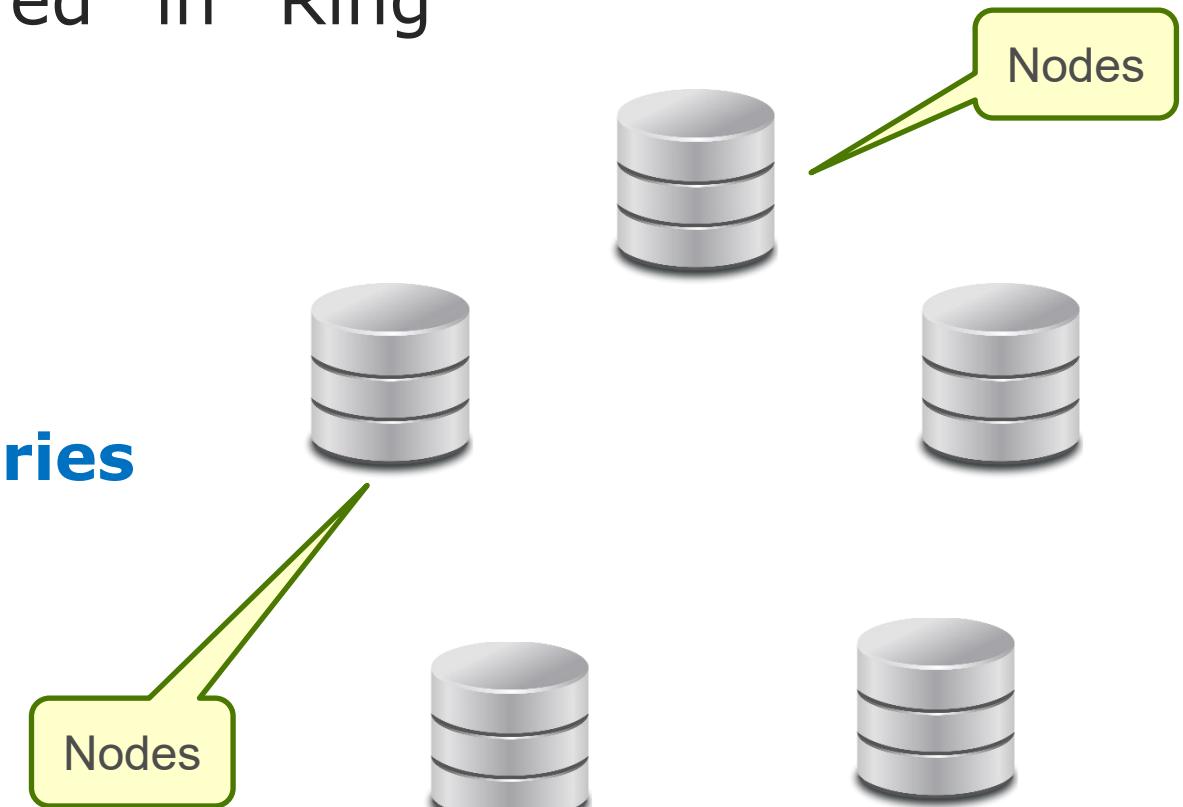
# When is Cassandra an option to use?

- You have a large amount of data that does not fit on one server
- The structure of the data changes often
- The system needs to be available all the time
  - No downtime
- Optimized read performance



# Cassandra: Distributed Partitioned DB

- Nodes are logically structured in Ring Topology.
- No joins
- Query should run on 1 node
  - **Partition data based on queries**
- Denormalization
  - Optimized reads
  - Data duplication





# Denormalization

- Denormalization is a strategy used on a database to increase performance.
- In computing, denormalization is the process of trying to improve the read performance of a DB, at the expense of losing some write performance, by adding redundant copies of data.
- Pros: Quick Read, simple queries
- Cons: Multiple writes, Manual integrity



# Partitioning – Data is Distributed

customerNr	firstName	lastName
1	Frank	Brown
2	Bob	Johnson
3	John	Jackson
4	Frank	Young
5	Sue	Jones



customerNr	firstName	lastName
1	Frank	Brown
2	Bob	Johnson

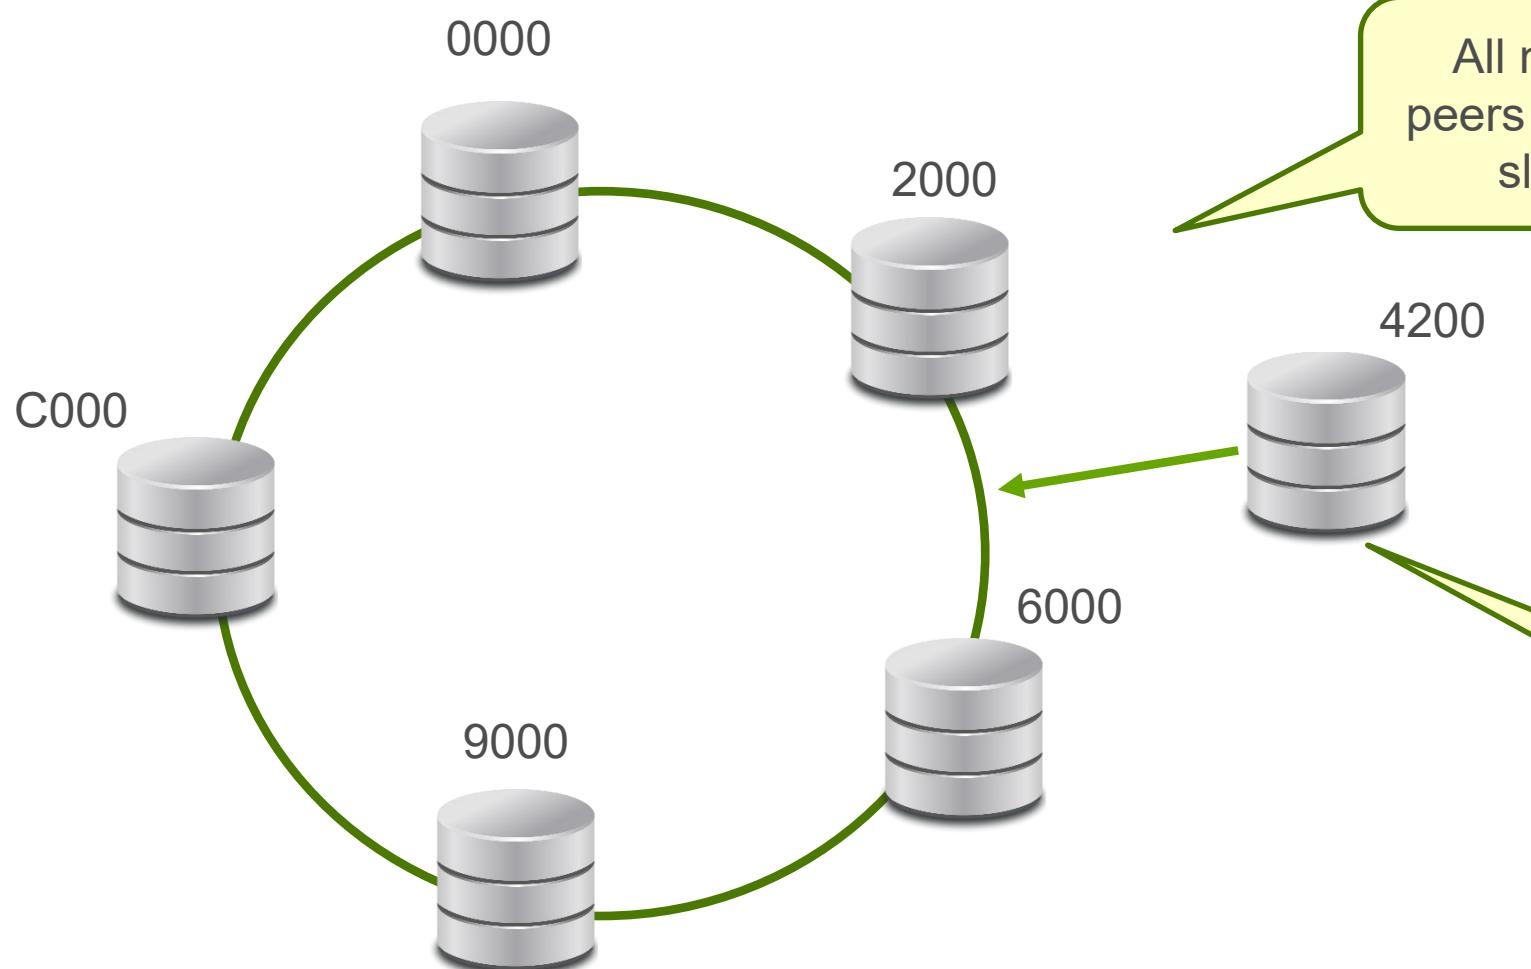
customerNr	firstName	lastName
3	John	Jackson
5	Sue	Jones

customerNr	firstName	lastName
4	Frank	Young

Hashed value of key associated with data partition is used to assign it to a node in the ring.



# Cassandra Architecture

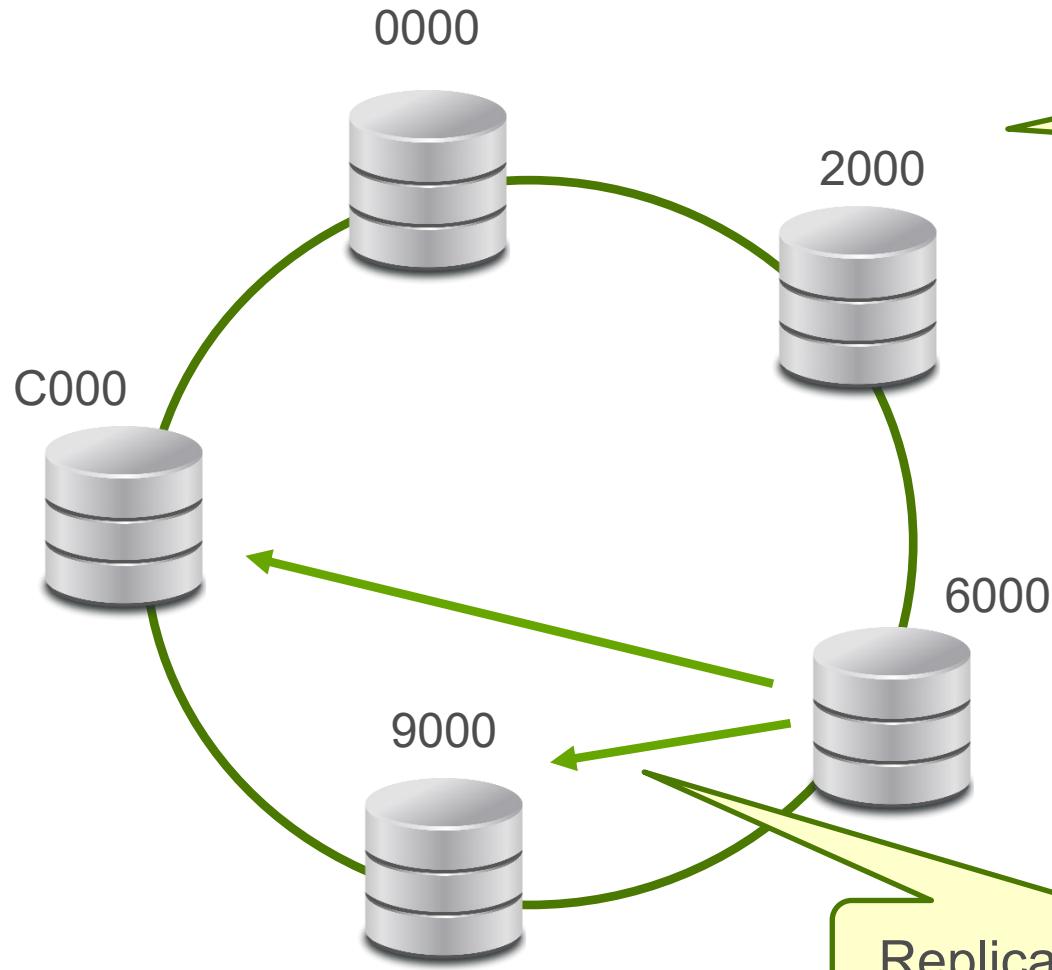


All nodes are peers (no master, slave,...)

Nodes can be added and removed without downtime



# Data Replication for Fault Tolerance



Consistency can be tuned for read and write. You can choose consistency level: ALL, Quorum, ONE

Consistency level typically represents a fraction of nodes within the cluster, which are required to acknowledge the operation for it to be considered successful.

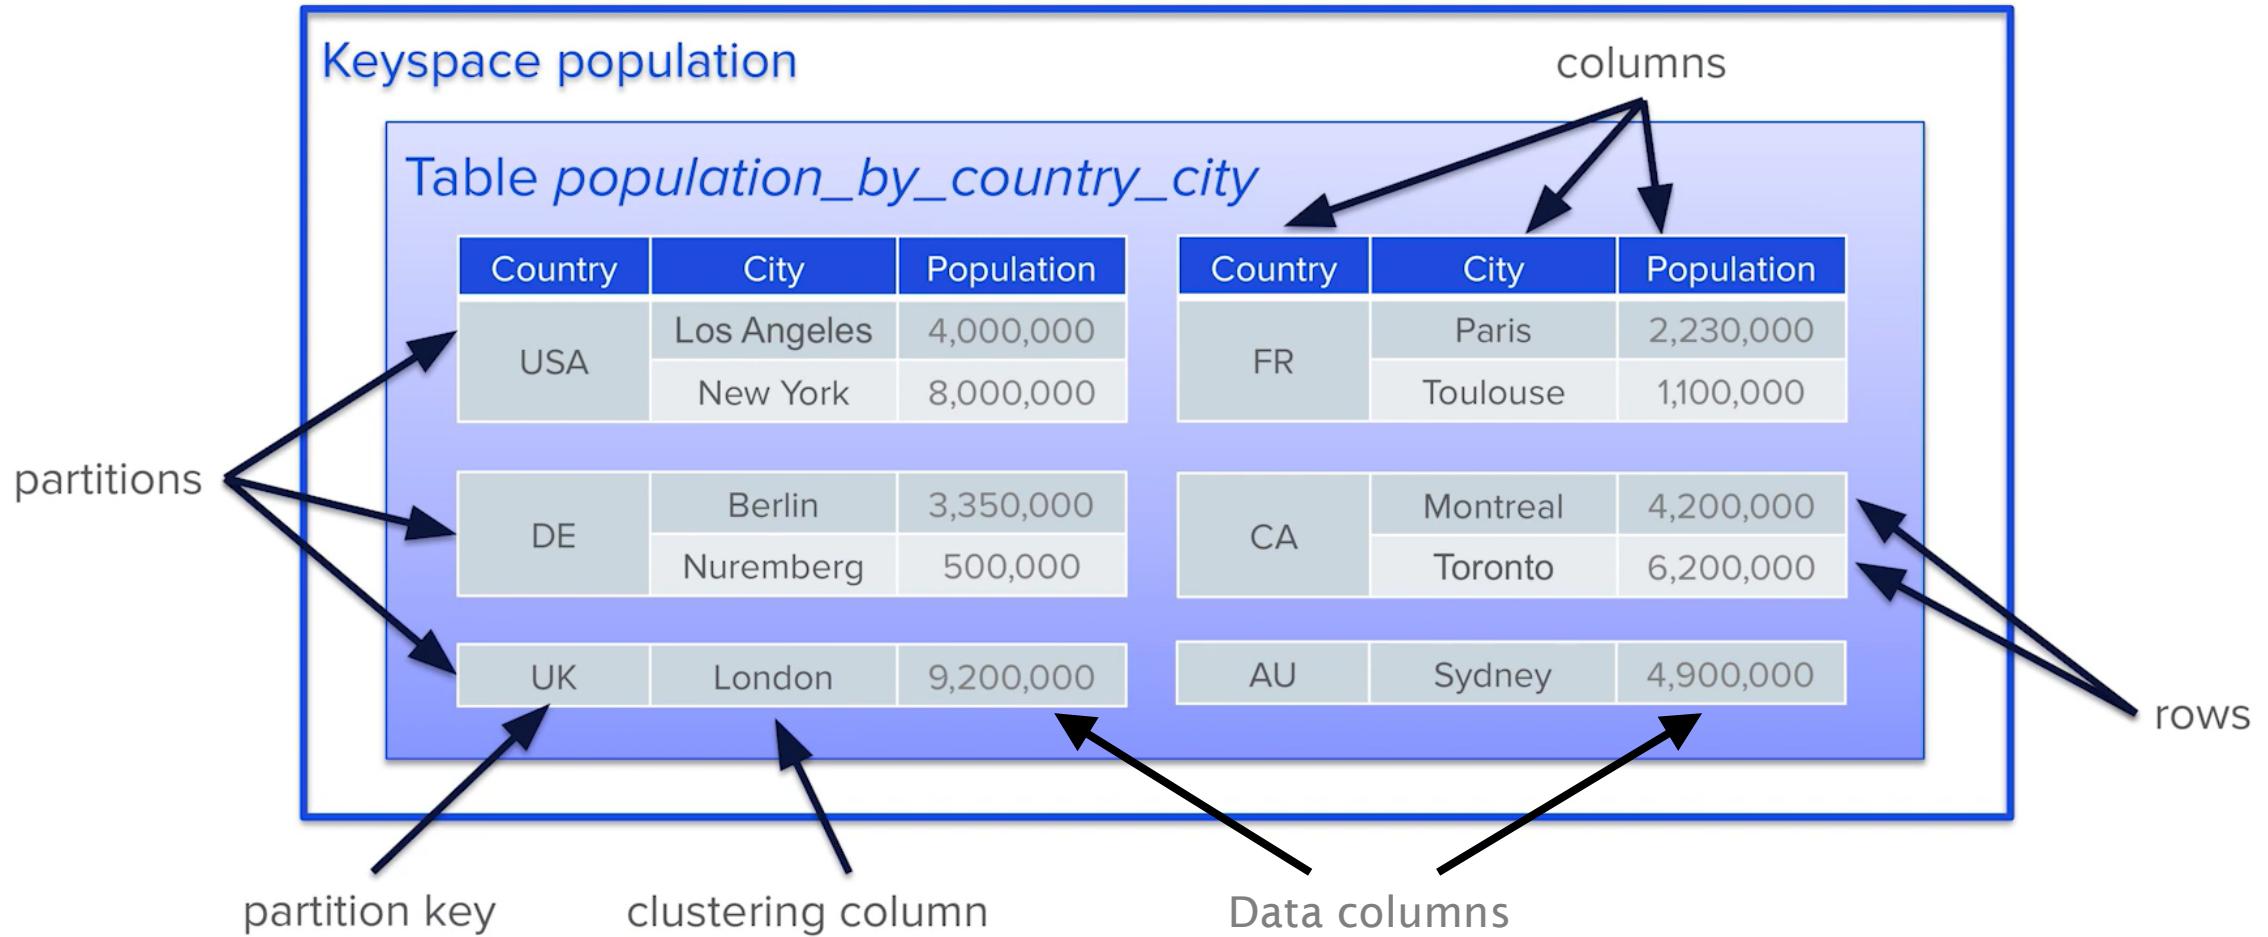


# Quorum Consistency Level

- **Quorum** consistency helps in preventing inconsistencies and conflicts in a distributed database.
  - For example, if you have a replication factor of 3 (meaning same data is stored on 3 nodes), you can use a quorum consistency level of QUORUM, which requires acknowledgment from at least 2 replicas for both reads and writes.
  - This means that you achieve a majority vote among the replicas for a read or write operation, which ensures strong consistency.
- There are more consistency levels other than ALL, Quorum and ONE



# Cassandra Data Model





## Cassandra Data Model contd..

- Tabular data model where data is organized in rows and columns.
- Keyspace is a container of tables. Like a database or schema in RDBMS but is more flexible/dynamic.
- Groups of related rows called *partitions* are stored together on the same node (or nodes).
- Each row contains a partition key.



# Creating a Table in CQL (Cassandra Query Language)

```
CREATE TABLE my_keyspace.my_table (
    partition_key_column_1 data_type,
    partition_key_column_2 data_type,
    clustering_column data_type,
    other_columns data_type,
    PRIMARY KEY ((partition_key_column_1,
                  partition_key_column_2), clustering_column)
);
```

Subset  
of SQL



# Creating a Table in CQL



Clustering  
columns naturally  
order the data.



# Primary Key

## Primary Key

- An identifier for a row. Consists of
  - at least one **Partition Key**
  - and zero or more **Clustering Columns**
- Must ensure uniqueness
- May define sorting through clustering columns

PRIMARY KEY ((A, B), C, D)

The diagram shows the primary key definition 'PRIMARY KEY ((A, B), C, D)' with annotations. Red arrows point from the labels 'Partition Key' to the columns A and B. Blue arrows point from the label 'Clustering Columns' to the columns C and D.

### Good Examples:

```
PRIMARY KEY ((city), last_name, first_name, email);
```

```
PRIMARY KEY (user_id);
```

### Bad Example:

```
PRIMARY KEY ((city), last_name, first_name);
```





# Partition key

- An identifier for a partition which **partitions rows**.
- Decides where to place this record
- Consists of at least one column, may have more if needed
- Needs to be mentioned while querying the table

```
CREATE TABLE killrvideo.users_by_city (
    city text,
    last_name text,
    first_name text,
    address text,
    email text,
    PRIMARY KEY ((city), last_name, first_name, email));
```

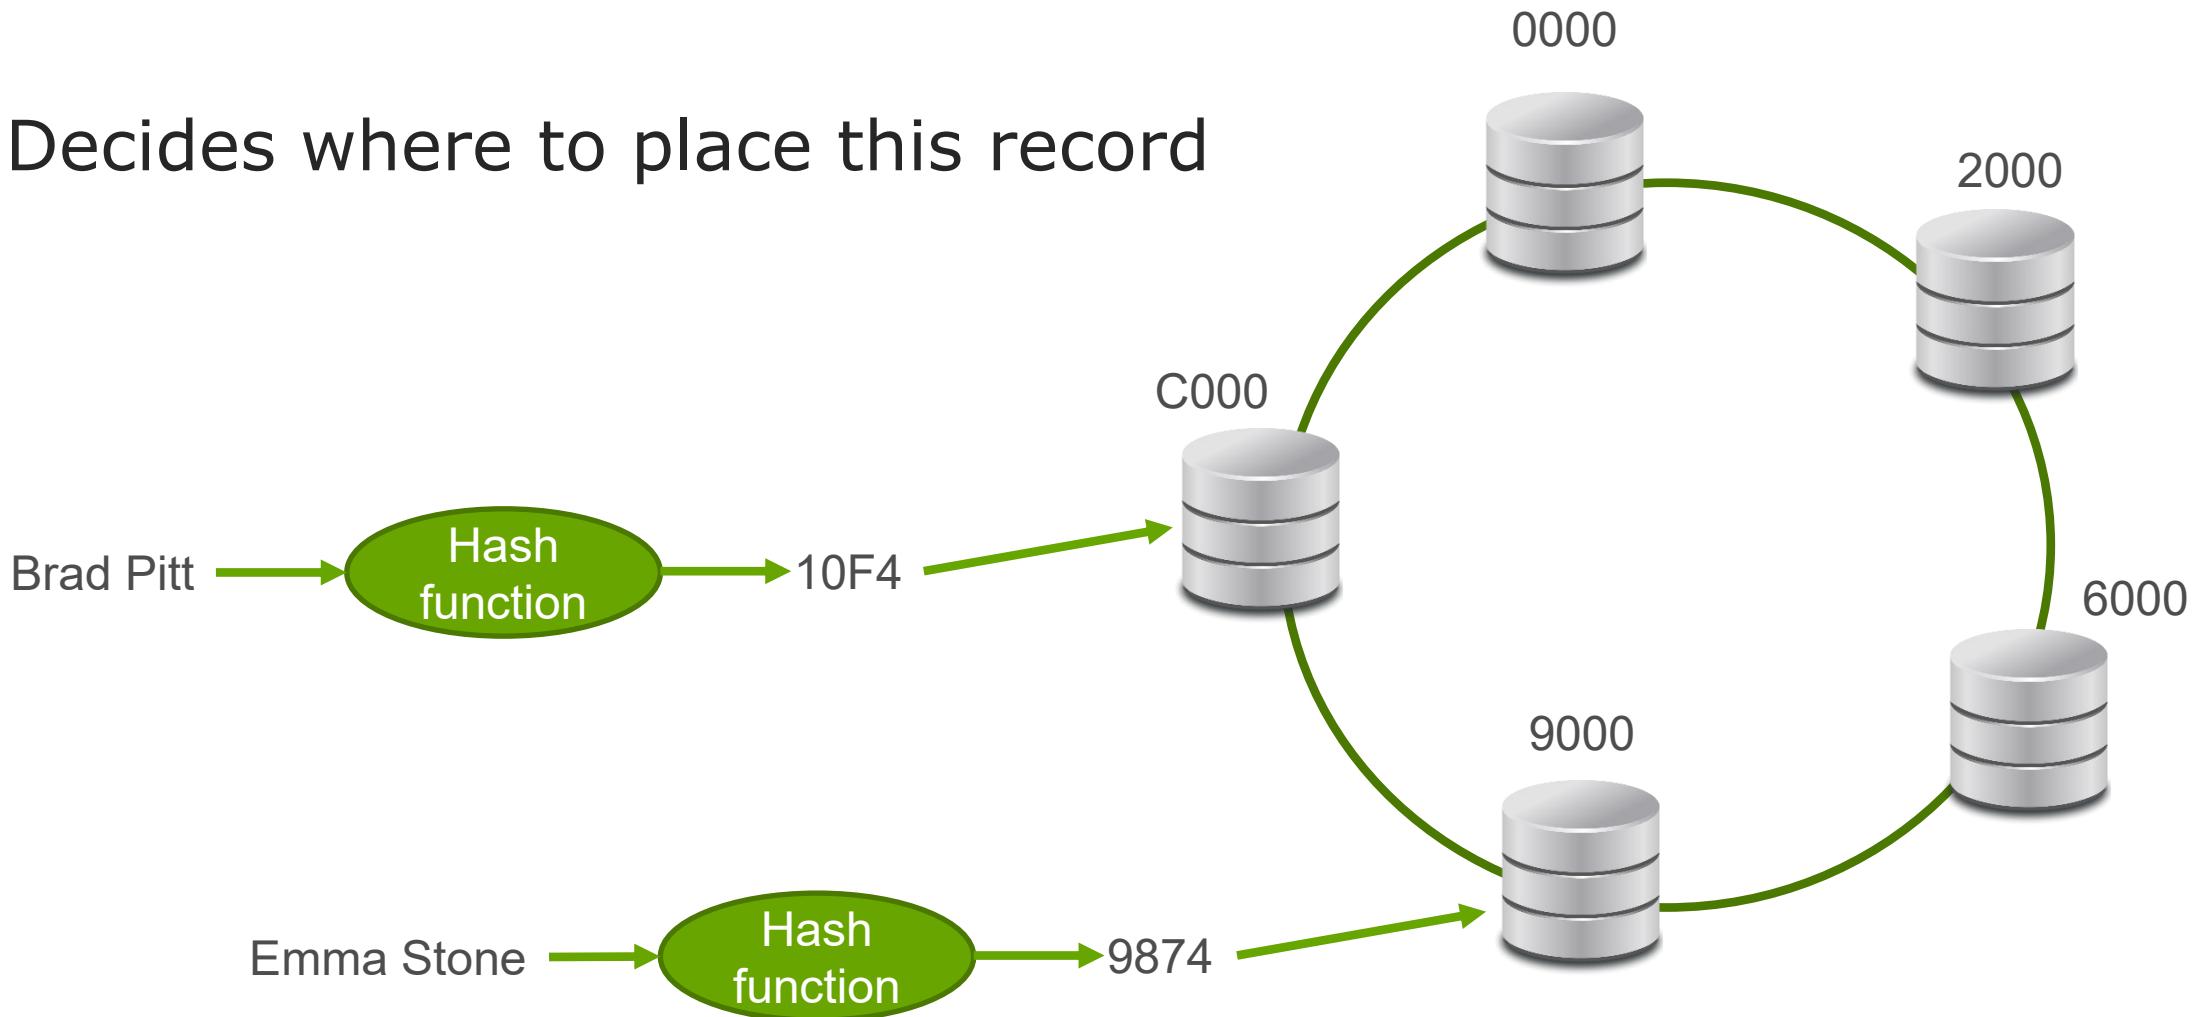
Partition key

Clustering columns



# Partition key

- Decides where to place this record





# Clustering Columns

- Used either to ensure uniqueness or sorting order for rows.
- Group rows together so they can be found fast in a query
- Attributes you want to query on besides the partition key
- Optional to add

```
CREATE TABLE killrvideo.users_by_city (
    city text,
    last_name text,
    first_name text,
    address text,
    email text,
    PRIMARY KEY ((city), last_name, first_name, email));
```

Partition key

Clustering columns

PRIMARY KEY ((city), last\_name, first\_name);

**Not Unique**

PRIMARY KEY ((city), last\_name, first\_name, email);

PRIMARY KEY ((video\_id), comment\_id);

**Not Sorted**

PRIMARY KEY ((video\_id), created\_at, comment\_id);



# Primary Key cannot be changed!

- Once you create a table with a Primary Key, you cannot change that Primary Key. You can add or remove as many data columns as you want but you cannot change the Primary Key.
- Why?
- Because, the Primary Key is defining the physical location of the data around the cluster (which row should go to which node), so if you change that, you will need to change the physical location of all the data.



# Rules of a Good Partition

## 1. Store together what you retrieve together

Example: open a video? Get the comments in a single query!

PRIMARY KEY ((video\_id), created\_at, comment\_id); 

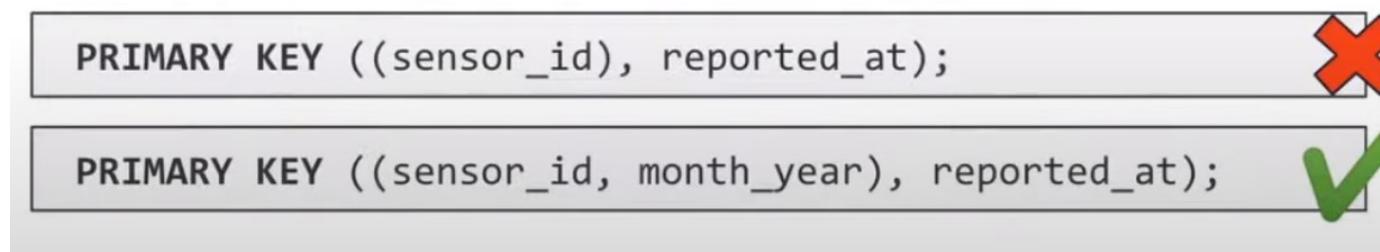
PRIMARY KEY ((comment\_id), created\_at); 



## Rules of a Good Partition contd..

### 2. Avoid big and constantly growing partitions

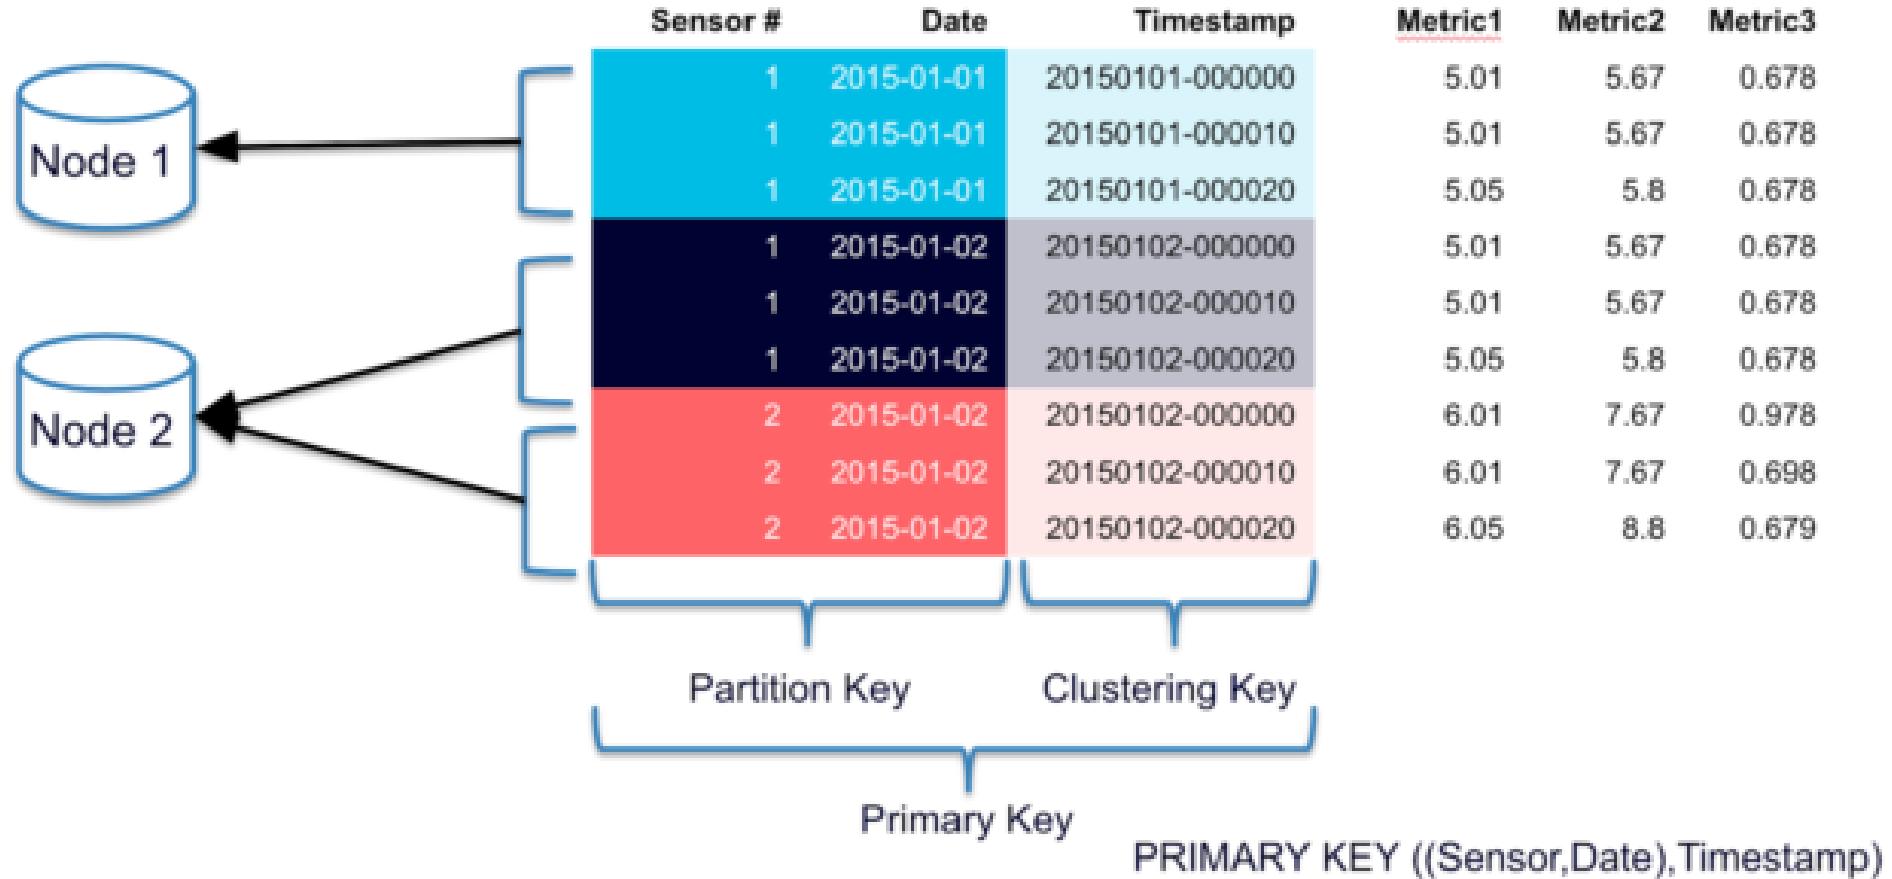
- Big partitions will increase read times because you'll need to query multiple servers.
- Up to 2 billion cells per partition, ~100K rows or ~100MB is ideal



# BUCKETING



# Bucketing Example

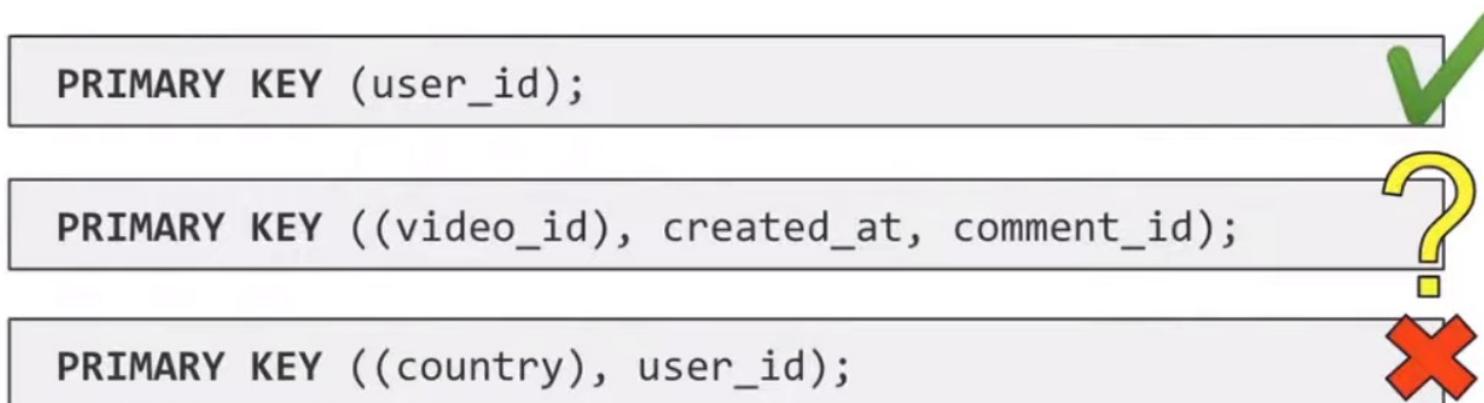




# Rules of a Good Partition contd..

## 3. Avoid hot partitions

- Hot partition is a partition that is getting hit a lot more frequently compared to other partitions.
- Good data modelling is spreading the data across all the partitions evenly so that all the nodes are hit at an equal rate.





# Data types

- TEXT, INT, BIGINT, FLOAT, DOUBLE, BOOLEAN, DECIMAL
- UUID
- TIMEUUID
- Collection
  - SET
  - LIST
  - MAP

Always use UUID (Universally Unique Identifier) for generated unique key

Always use TIMEUUID for generated unique key alongwith a timestamp

```
cqlsh> Create table University.Teacher  
...  
... id int,  
... Name text,  
... Email set<text>,  
... Primary key(id)  
...);  
cqlsh>
```

collection of emails

```
cqlsh> insert into University.Teacher(id,Name,Email) values(1,'Guru99',{'abc@gmail.com','xyz@hotmail.com'});  
cqlsh>
```

insertion in collection



# Naming Convention

- Snake\_case is mostly used
- You can add the partition column name after “by” to make the table name more meaningful.
  - users\_by\_rating
  - population\_by\_country



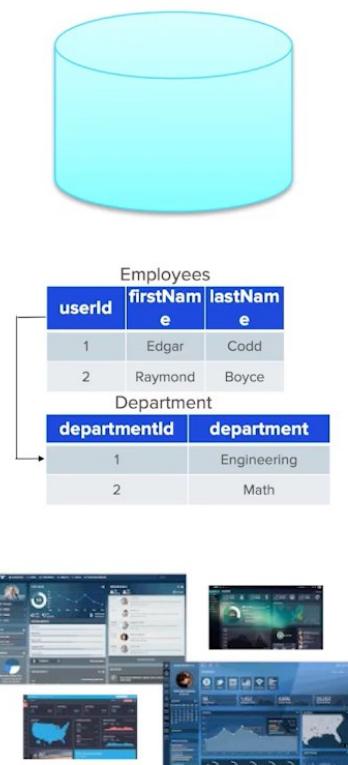
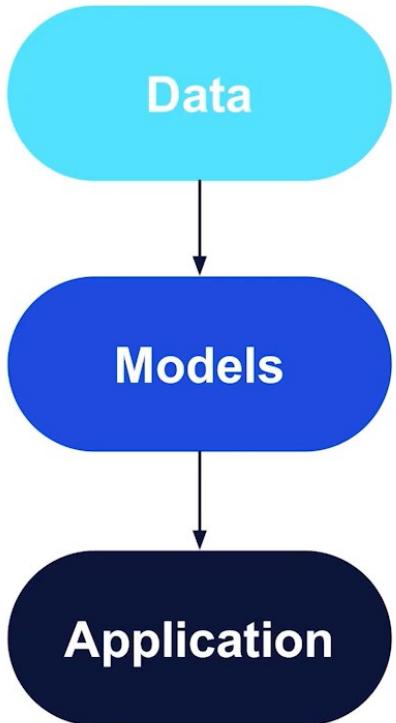
# Data modeling

- When migrating a relational database to Cassandra, it's important to understand that Cassandra is a NoSQL database designed for **scalability** and **high availability**. As a result, the data modeling approach is quite different from that of a traditional relational database.
- In Cassandra, you need to design your tables based on your specific query patterns.
- This approach is known as "**query-driven data modeling**" in Cassandra.

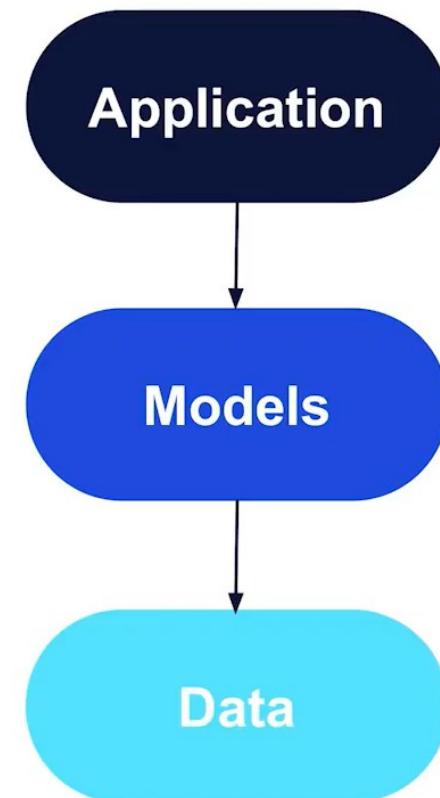


# Data modeling

- RDBMS



- Cassandra



id	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math





# Data modeling

## **Relational data modelling:**

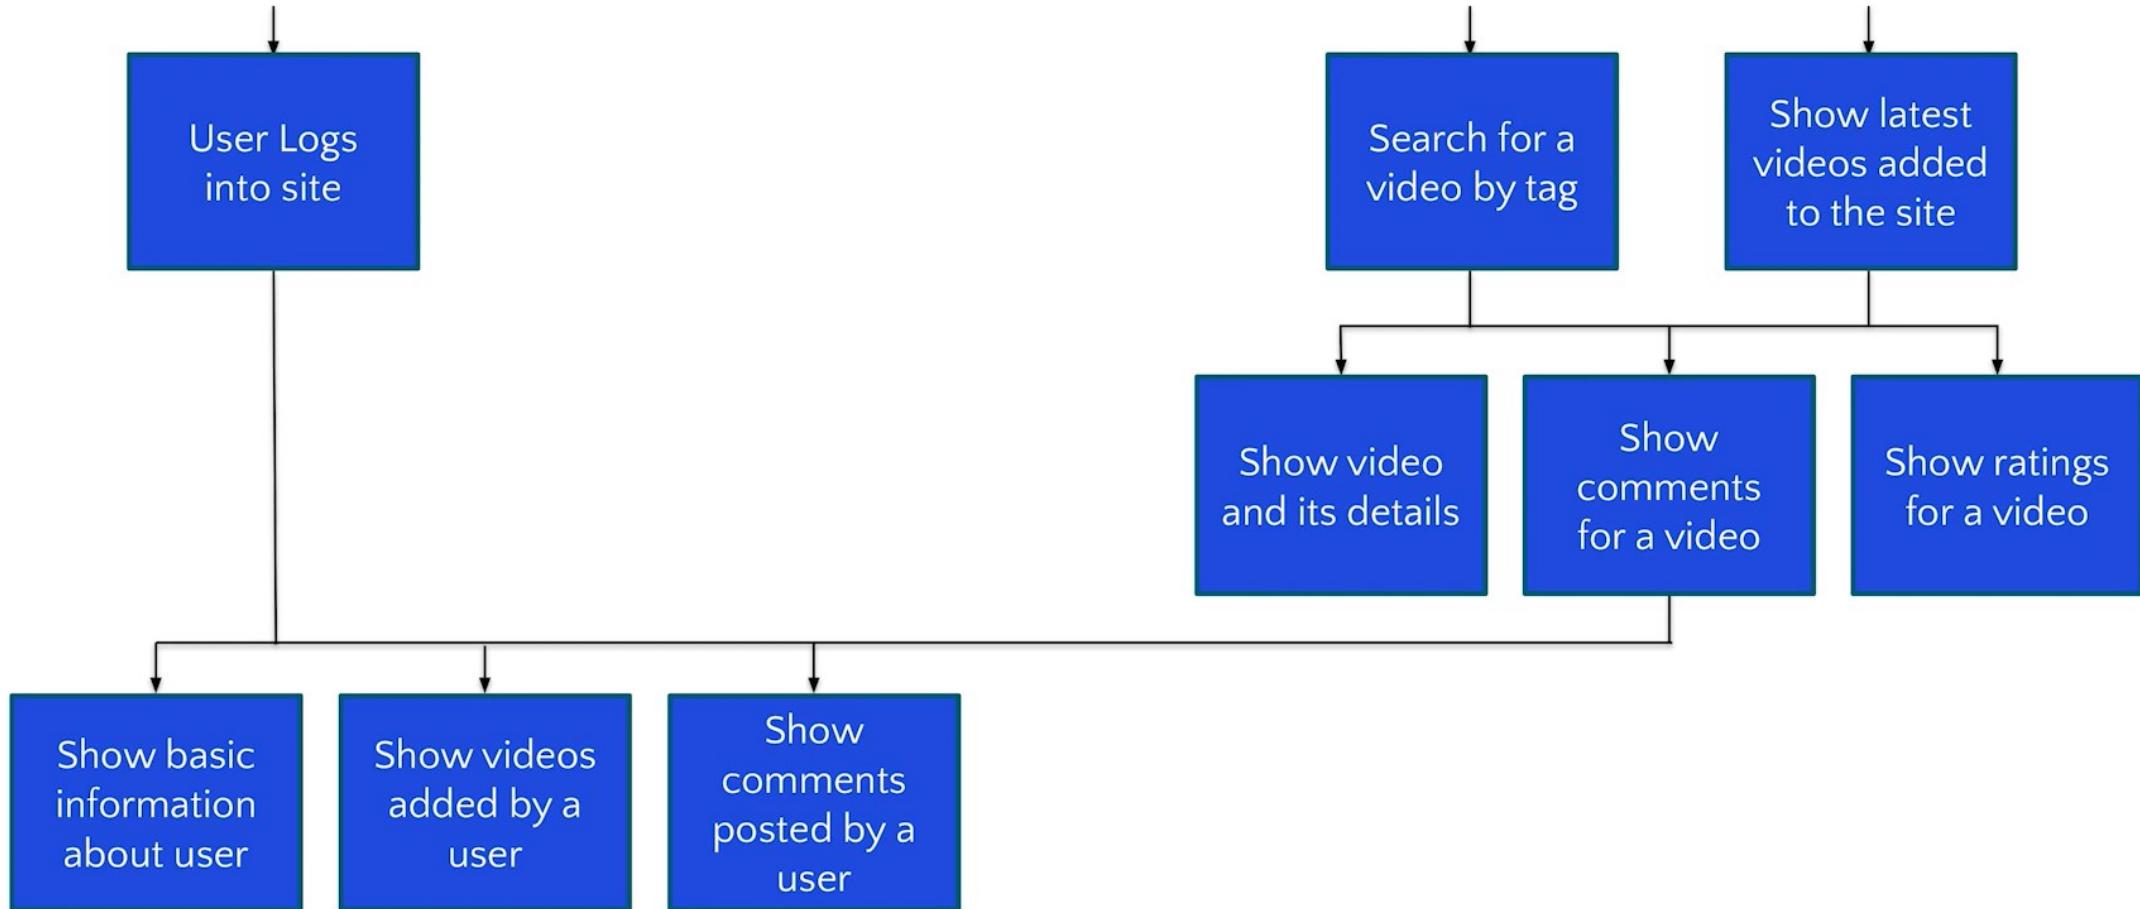
1. Analyze raw data
2. Identify entities, their properties and relations
3. Design tables, using normalization and FKs.
4. Use JOIN when doing queries to join normalized data from multiple tables.

## **NoSQL Data Modelling:**

1. Analyze user behavior (customer first!)
2. Identify workflows, their dependencies and needs
3. Define Queries to fulfill these workflows
4. Knowing the queries, design tables, using denormalization.
5. Use BATCH when inserting or updating denormalized data of multiple tables.



# Video Application Workflow





# Queries that support the workflow

## Users

User Logs  
into site

Find user by email  
address

Show basic  
information  
about user

Find user by id

## Comments

Show  
comments  
for a video

Find comments by  
video (latest first)

Show  
comments  
posted by a  
user

Find comments by  
user (latest first)

## Ratings

Show ratings  
for a video

Find ratings by video

## Videos

Search for a  
video by tag

Find video by tag

Show latest  
videos added  
to the site

Find videos by date  
(latest first)

Show video  
and its details

Find video by id

Show videos  
added by a  
user

Find videos by user  
(latest first)



# Relational way

User Logs  
into site

Find user by email  
address

Show basic  
information  
about user

Find user by id

- Single Users table with all user data and PK as userID.

```
CREATE TABLE users (
   (userID int,
    firstName text,
    lastName text,
    email text,
    createdDate timestamp,
    PRIMARY KEY (userID)
);
```



# Cassandra way

User Logs  
into site

Find user by email  
address

Show basic  
information  
about user

Find user by id

```
CREATE TABLE users_by_email (
    email text,
    password text,
    userid uuid,
    PRIMARY KEY (email)
);
```

```
CREATE TABLE users_by_id (
    userid uuid,
    firstname text,
    lastname text,
    email text,
    created_date timestamp,
    PRIMARY KEY (userid)
);
```

In the Cassandra way, you'd have multiple tables optimized for specific query patterns, and data duplication is common to support various access paths efficiently.



# Video queries

Show video  
and its details

Find video by id

Show videos  
added by a  
user

Find videos by user  
(latest first)

```
CREATE TABLE video_by_id (
    videoid uuid,
    userid uuid,
    name text,
    description text,
    location text,
    location_type int,
    preview_image_location text,
    tags set<text>,
    added_date timestamp,
    PRIMARY KEY (videoid)
);
```

```
CREATE TABLE videos_by_user (
    userid uuid,
    added_date timestamp,
    videoid uuid,
    name text,
    preview_image_location text,
    PRIMARY KEY ((userid), added_date, videoid)
)
WITH CLUSTERING ORDER BY (
    added_date DESC,
    videoid ASC);
```



# CQL Cassandra Query Language: Create Table

**ratings\_by\_user**

email	title	year	rating
joe@datastax.com	Alice in Wonderland	2010	9
joe@datastax.com	Edward Scissorhands	1990	10
jen@datastax.com	Alice in Wonderland	1951	8
jen@datastax.com	Alice in Wonderland	2010	10

```
CREATE TABLE ratings_by_user (
    email TEXT,
    title TEXT,
    year INT,
    rating INT,
    PRIMARY KEY ((email), title, year)
); ↵
```

Partition key

Clustering  
columns



# CQL Cassandra Query Language: Insert

email	title	year	rating
joe@datastax.com	Alice in Wonderland	2010	9
joe@datastax.com	Edward Scissorhands	1990	10
jen@datastax.com	Alice in Wonderland	1951	8
jen@datastax.com	Alice in Wonderland	2010	10

```
CREATE TABLE ratings_by_user (
    email TEXT,
    title TEXT,
    year INT,
    rating INT,
    PRIMARY KEY ((email), title, year)
); ↵
```

```
INSERT INTO ratings_by_user (email, title, year,
rating)
VALUES ('joe@datastax.com', 'Alice in Wonderland',
2010, 9);
INSERT INTO ratings_by_user (email, title, year,
rating)
VALUES ('joe@datastax.com', 'Edward Scissorhands',
1990, 10);
INSERT INTO ratings_by_user (email, title, year,
rating)
VALUES ('jen@datastax.com', 'Alice in Wonderland',
2010, 10);
INSERT INTO ratings_by_user (email, title, year,
rating)
VALUES ('jen@datastax.com', 'Alice in Wonderland',
1951, 8); ↵
```



# Ratings by movie

title	year	email	rating
Alice in Wonderland	2010	jen@datastax.com	10
Alice in Wonderland	2010	joe@datastax.com	9
Alice in Wonderland	1951	jen@datastax.com	8
Edward Scissorhands	1990	joe@datastax.com	10

```
CREATE TABLE ratings_by_movie (
    title TEXT,
    year INT,
    email TEXT,
    rating INT,
    PRIMARY KEY ((title, year), email)
); ↵
```

```
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Alice in Wonderland', 2010,
'jen@datastax.com', 10);
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Alice in Wonderland', 2010,
'joe@datastax.com', 9);
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Alice in Wonderland', 1951,
'jen@datastax.com', 8);
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Edward Scissorhands', 1990,
'joe@datastax.com', 10); ↵
```



# Actors by movie

title	year	first_name	last_name
Alice in Wonderland	2010	Anne	Hathaway
Alice in Wonderland	2010	Johnny	Depp
Edward Scissorhands	1990	Johnny	Depp

```
CREATE TABLE actors_by_movie (
    title TEXT,
    year INT,
    first_name TEXT,
    last_name TEXT,
    PRIMARY KEY ((title, year), first_name, last_name)
); ↵
```



# Movies by actor

first_name	last_name	title	year
Johnny	Depp	Alice in Wonderland	2010
Johnny	Depp	Edward Scissorhands	1990
Anne	Hathaway	Alice in Wonderland	2010

```
CREATE TABLE movies_by_actor (
    first_name TEXT,
    last_name TEXT,
    title TEXT,
    year INT,
    PRIMARY KEY ((first_name, last_name), title, year)
); ↵
```



# Movies by user

email	watched_on	title	year
joe@datastax.com	2020-04-28	Edward Scissorhands	1990
joe@datastax.com	2020-03-08	Alice in Wonderland	2010
joe@datastax.com	2020-02-13	Toy Story 3	2010
joe@datastax.com	2020-01-22	Despicable Me	2010
joe@datastax.com	2019-12-30	Alice in Wonderland	1951
jen@datastax.com	2011-10-01	Alice in Wonderland	2010

```
CREATE TABLE movies_by_user (
    email TEXT,
    title TEXT,
    year INT,
    watched_on DATE,
    PRIMARY KEY ((email), watched_on, title, year)
) WITH CLUSTERING ORDER BY (watched_on DESC); ↵
```



# Data duplication

email	title	year	rating
joe@datastax.com	Alice in Wonderland	2010	9
joe@datastax.com	Edward Scissorhands	1990	10
jen@datastax.com	Alice in Wonderland	1951	8
jen@datastax.com	Alice in Wonderland	2010	10

title	year	first_name	last_name
Alice in Wonderland	2010	Anne	Hathaway
Alice in Wonderland	2010	Johnny	Depp
Edward Scissorhands	1990	Johnny	Depp

email	watched_on	title	year
joe@datastax.com	2020-04-28	Edward Scissorhands	1990
joe@datastax.com	2020-03-08	Alice in Wonderland	2010
joe@datastax.com	2020-02-13	Toy Story 3	2010
joe@datastax.com	2020-01-22	Despicable Me	2010
joe@datastax.com	2019-12-30	Alice in Wonderland	1951
jen@datastax.com	2011-10-01	Alice in Wonderland	2010

title	year	email	rating
Alice in Wonderland	2010	jen@datastax.com	10
Alice in Wonderland	2010	joe@datastax.com	9
Alice in Wonderland	1951	jen@datastax.com	8
Edward Scissorhands	1990	joe@datastax.com	10



# Queries

Q1. Retrieve all rows:

```
SELECT * FROM users; ↵
```

Never do this type of a query on real production data in Cassandra!

Q2. Retrieve one row/partition:

```
SELECT * FROM users  
WHERE email = 'joe@datastax.com'; ↵
```

Q3. Retrieve two rows/partitions:

```
SELECT * FROM users  
WHERE email IN ('joe@datastax.com',  
                 'jen@datastax.com'); ↵
```



# Queries

title	year	email	rating
Alice in Wonderland	2010	jen@datastax.com	10
Alice in Wonderland	2010	joe@datastax.com	9
Alice in Wonderland	1951	jen@datastax.com	8
Edward Scissorhands	1990	joe@datastax.com	10

A query can only filter on fields in the primary key

```
CREATE TABLE ratings_by_movie (
    title TEXT,
    year INT,
    email TEXT,
    rating INT,
    PRIMARY KEY ((title, year), email)
); ↵
```

You cannot directly filter on rating



# Queries

```
SELECT COUNT(rating) AS count,  
       SUM(rating) AS sum,  
       AVG(CAST(rating AS FLOAT)) AS avg,  
       MIN(rating) AS min,  
       MAX(rating) AS max  
FROM   ratings_by_movie  
WHERE  title = 'Alice in Wonderland'  
AND    year  = 2010; ↵
```

```
SELECT * FROM ratings_by_user  
WHERE email = 'jim@datastax.com'  
ORDER BY title ASC, year DESC;
```

```
SELECT  title, year,  
        AVG(CAST(rating AS FLOAT)) AS avg_rating  
FROM    ratings_by_movie  
GROUP BY title, year; ↵
```